

Consistency - Chapter 5

- Introduce several notions of Local Consistency:
 - arc consistency,
 - hyper-arc consistency,
 - k-consistency and strong k-consistency.
- Local consistency is about the existence of partial solutions and their extensions.
- Local consistency helps search by making partial solutions easier to find.

Example

Consider the three integer variables x , y and z each with the domain $\{0, 1, \dots, 10\}$ and the single constraint:

$$x + y = z$$

For any variable and any value there is a solution containing that value and variable.

Example Continued

$x, y, z \in \{0, 1, \dots, 10\}$ and $x + y = z$.

- For any value, x , of X set Y to be 0 and Z to be x ;
- For any value, y , of Y set X to be 0 and Z to be y ;
- For any value, z , of Z set X to be 0 and Y to be z .

But if we set X to be 8 then this imposes the restrictions:

- $0 \leq Y \leq 2$
- $8 \leq Z \leq 10$.

Local Consistency

- The various notions of local consistency help understand what is going on in the previous example.
- Constraint propagation attempts to reduce domains so that in the previous example assigning 8 to X reduces the domain of Y and Z so that during search fewer possibilities have to be explored.
- Propagation is the heart of modern constraint programming. Without it constraint programming just reduces to generate and test.

Arc Consistency of a Binary Constraint

- A binary constraint C on the variables x and y with domains X and Y is a subset of $X \times Y$. Such a $C \subseteq X \times Y$ is *arc consistent* if:
 - $\forall a \in X$ there $\exists b \in Y$ such that $(a, b) \in C$;
 - $\forall b \in Y$ there $\exists a \in X$ such that $(a, b) \in C$.
- Note both directions are important.
- A CSP is said to be *consistent* if all its binary constraints are consistent (this definition is unproblematic if all the constraints in the CSP are binary).

Examples of Arc-consistency

- For $x \in [2 \dots 6]$, $y \in [3 \dots 7]$ the constraint

$$x < y$$

is arc-consistent.

- For example
 - if $x = 2$ then there is a solution
 - if $x = 6$ then y must be 7
 - if $y = 3$ then x must be 2.

A Non Arc-Consistent Constraint

- For $x \in [2..7]$ and $y \in [3..7]$ the constraint:

$$x < y$$

is not arc consistent.

- If $x = 7$ (allowed by the domains) then there is no value for y satisfying the constraint.

Status of Arc Consistency

- Arc consistency does not imply consistency. Given $x \in \{a, b\}$ and $y \in \{a, b\}$ the two constraints

$$x = y \text{ and } x \neq y$$


are both arc-consistent.

- but there is obviously no solution to *both* constraints.

Status of Arc Consistency

For particular CSPs arc consistency implies consistency.

- Given a CSP



where each constraint is arc-consistent, the whole CSP is consistent.

- To see this pick a value for y then arc-consistency gives a value for x and z .

In general if the constraint graph is a tree then arc consistency implies consistency.

Achieving and Using Arc-Consistency

- The Basic Idea
 - remove values from the domain which do not take part in solutions.
- Given a constraint $C \subseteq D_x \times D_y$ on the variables $x \in D_x$ and $y \in D_y$ reduce the domains by the following two rules:
 - $D'_x = \{a \in D_x \mid \exists b \in D_y \text{ s.t. } (a, b) \in C\}$
 - $D'_y = \{b \in D_y \mid \exists a \in D_x \text{ s.t. } (a, b) \in C\}$
- To achieve arc consistency you must apply both rules.

Directional Arc Consistency

- Sometimes it can be expensive (computationally) to achieve arc consistency: perhaps because the domains are large.
- While searching for values of variables with backtrack search you might not need full arc consistency.
- Suppose you always assign a value to x before you assign a value to y then given a constraint C on x and y all you need is that for all values, a , of x there is a tuple (a, b) in C giving a value for y .

This leads to the notion of directional consistency.

Directional Arc Consistency

Ingredients:

- A linear order \preceq on the variables.
- A linear order is a binary relation such that:
 - For all x , $x \preceq x$ (reflexivity)
 - For all x and y , $x \preceq y$ and $y \preceq x$ implies $x = y$ (antisymmetry)
 - For all x, y and z , $x \preceq y$ and $y \preceq z$ implies $x \preceq z$ (transitivity)
 - For all x and y either $x \preceq y$ or $y \preceq x$.
- A linear order is just some fixed order on the variables.

Directional Arc Consistency

Assume a linear ordering \preceq on the variables:

- A constraint C on $x \in D_x, y \in D_y$ is directionally arc consistent w.r.t. \preceq if:
 - $\forall a \in D_x$ there $\exists b \in D_y$ such that $(a, b) \in C$ provided that $x \preceq y$.
 - $\forall b \in D_y$ there $\exists a \in D_x$ such that $(a, b) \in C$ provided that $y \preceq x$.
 - A CSP is directionally arc consistent w.r.t. \preceq if all its binary constraints are.

Directionally Arc Constraints - Example

- Given $x \in [2 \dots 7]$, $y \in [3 \dots 7]$ the constraint $x < y$ is not arc consistent (no solution for $x = 7$).
- It is directionally arc consistent w.r.t. $y \preceq x$. That is for any value you assign to y there is an assignment to x satisfying the constraint.
- Is is not directionally consistent when $x \preceq y$: for example assigning 7 to x means that we can not assign any value to y .

Non Binary Constraints and Consistency

- Most constraints that you meet are not binary, for example you have already met the `alldifferent` constraint.
- Although you can always model a problem with binary constraints it is not always as efficient as using global (non-binary) constraints.
- There are lots of possible definitions of arc-consistency possible for non-binary constraints.
- We will look at a pragmatically useful one which is often used in implementations.

Hyper-Arc Consistency

- A constraint on the variables x_1, \dots, x_n with the domains D_1, \dots, D_n is *hyper-arc consistent* if

$$\forall i \in [1..n] \forall a \in D_i \text{ there } \exists d \in C \text{ s.t. } a = d[x_i]$$

- The notation $d[x_i]$ takes a tuple of values in a relation and projects it down to the entry corresponding to the variable x_i .
- The tuple d is often called a supporting tuple of the assignment $x = a$.

Hyper-arc Consistency - Examples

Suppose C is a constraint on the variables $x \in \{1, 2, 3\}$, $y \in \{1, 2, 3\}$ and $z \in \{1, 2, 3\}$.

Suppose $C(x, y, z)$ is given by a list of tuples: $\langle 1, 2, 3 \rangle, \langle 1, 2, 2 \rangle$ and $\langle 2, 3, 3 \rangle$ then this is not hyper-arc consistent w.r.t to the domains since for $z = 1$ there is no tuple supporting it.

The first constraint in this lecture is hyper-arc consistent ($x, y, z \in \{0, 1, \dots, 10\}$ and $x + y = z$).

Global Constraints

- Given some global constraints how to prune values from the domain to keep hyper-arc consistency.
- Problem to do this efficiently.
- Later on in the course you will meet many global constraints and pruning algorithms.

Instantiations

Fix a CSP \mathcal{P} .

- An instantiation is a function on a subset of the variables of \mathcal{P} which assigns a value in the domain.
- Notation from the book:

$$\{(x_1, d_1), \dots, (x_n, d_n)\}$$

means x_1 is assigned to d_1 , \dots , x_n is assigned to d_n .

- Another common notation:

$$x_1 \mapsto d_1 \wedge \dots \wedge x_n \mapsto d_n$$

Consistent Instantiations

We want a notion of when a partial solution satisfies a CSP \mathcal{P} .

- Given an instantiation I on the variables X , we denote the restriction of I to a set $Y \subset X$:

$$I|Y$$

- An instantiation I with domain X is *consistent* if for every constraint C of \mathcal{P} on the variables Y with $Y \subset X$, $I|Y$ satisfies C . (I will often refer to consistent instantiations as partial solutions).
- A Consistent instantiation is a *k-consistent* instantiation if its domain consists of k variables.

Example

Let \mathcal{P} be the CSP

$$x < y, y < z, x < z; x \in [0 \dots 4], y \in [1 \dots 5], z \in [5 \dots 10]$$

Let I be $x \mapsto 0 \wedge y \mapsto 5 \wedge z \mapsto 6$


- $I|_{\{x, y\}} = x \mapsto 0 \wedge y \mapsto 5$ and satisfies $x < y$;
- $I|_{\{x, z\}} = x \mapsto 0 \wedge z \mapsto 6$ and satisfies $x < z$;
- $I|_{\{y, z\}} = y \mapsto 5 \wedge z \mapsto 6$ and satisfies $y < z$.
- So I is a 3-consistent instantiation. It is a solution to the CSP.

k -Consistency

- A CSP is *1-consistent* if for every variable x every unary constraint equals the domain of x (Node consistency).
- A CSP is *k -consistent* for $k > 1$ if every $k - 1$ -consistent instantiation can be extended to a k -consistent instantiation *no matter* which new variable is chosen.

Example

Consider the CSP.



This is 2-consistent.


Pick any partial solution:

$$x_1 \mapsto 1$$

pick any other variable say x_3 and we can find a consistent instantiation satisfying the constraints say $x_3 = 2$.

Example Continued

But



is not 3-consistent.

To prove this we pick some 2-consistent partial solution:

$$x_1 \mapsto 1 \wedge x_2 \mapsto 2$$

this cannot be extended to a solution on 3 variables.

Recap

- To show k consistency you have to look at every consistent $k - 1$ solution and every other variable and show that the extension exists.
- To disprove consistency you only have to find one counter example.


Why is Consistency a Good Thing?

- If your CSP is k consistent you know that if you have assigned $k - 1$ variables you can assign the next variable in the search tree.

But k -consistency is not the whole story.

- k consistency does not imply $k - 1$ consistency (example next slide);
- If your CSP is k consistent you still have to find a $k - 1$ consistent partial solution.

Example



This CSP is consistent but not 3 consistent. There is a solution:

$$x_1 \mapsto 1 \wedge x_2 \mapsto 2 \wedge x_3 \mapsto 3$$

but the partial solution:

$$x_3 \mapsto 1 \wedge x_2 \mapsto 2$$

has no extension.

Strong k -Consistency

- A CSP is *strongly k -consistent* where $k \geq 1$ if it is i -consistent for every $i \in [1 \dots k]$.
- A CSP with non-empty domains on k variables which is k -consistent has a solution.
- Sometimes we can do better (later in the course) and show when we only need a certain amount of local consistency to achieve global consistency.
- In the next lecture we will see that it is possible to make a CSP k -consistent (or show it is not consistent) for any k . So one way of solving the CSP is to progressively increase the level of consistency until the CSP is solved.

Path Consistency

Path Consistency is a special case of k -consistency when $k = 2$ (plus some other conditions to be spelled out below).

Consider the CSP

$$x < y, y < z, z < x$$

with $x \in \{1 \dots 1000\}$, $y \in \{1 \dots 1000\}$ and $z \in \{1 \dots 1000\}$.

This can be shown to be inconsistent by using the arc-consistency domain reduction rules. Applying the rule to $x < y$ gives that $x \in \{1 \dots 999\}$ then applying the rule $z < x$ gives $z \in \{1 \dots 998\}$ and so on. Eventually you can show that the CSP is inconsistent.

Path Consistency

But the CSP

$$x < y, y < z, z < x$$

can be shown to be inconsistent: by combining the two constraints $x < y$ and $y < z$ allows you to deduce that $x < z$ contradicting the constraint $z < x$.

Operations on Binary Relations

Path consistency generalises the previous example to arbitrary binary constraints.

Given two binary relations R and S define the following operations:

- the *transpose* of R by

$$R^T = \{(b, a) \mid (a, b) \in R\}$$

some people write R^{op} instead of R^T ;

- the composition of R and S by

$$R \cdot S = \{(a, c) \mid \exists b \text{ s.t. } (a, b) \in R \wedge (b, c) \in S\}$$

Examples of Composition

- Let

$$R = \{(1, 2), (1, 3), (4, 3), (2, 3)\}$$

- and

$$S = \{(1, 3), (3, 1), (4, 2)\}$$

- then

$$R \cdot S = \{(1, 1), (4, 1), (2, 1)\}$$

- and

$$S \cdot R = \{(3, 2), (3, 3), (4, 3)\}$$

Normalised CSPs

A CSP is *normalised* if for every subsequence x,y of variables at most one constraint on x,y exists (defn. 5.15)

The CSP

$$x + y < 5, \ x + y \neq 2, \ x \in \{0 \dots 4\}, \ y \in \{0 \dots 4\}$$

is not normalised.

A CSP can be made normalised by taking the conjunctions of the multiple constraints.

Path Consistency


Notation a C constraint on the variables x, y will be denoted $C_{x,y}$.

Given a constraint on the variables x, y there is also an imaginary constraint $C_{y,x}^T$ on the variables y, x which is the transpose of the constraint C .

A normalised CSP is *path consistent* (Defn. 5.18) if for every subset of variables $\{x, y, z\}$ of its variables have

$$C_{x,z} \subseteq C_{x,y} \cdot C_{y,z}$$

Example



Path Consistency

- To make a CSP arc-consistent we reduced the domains.
- To make it path-consistent we reduce the constraints.
- Path consistency is about triangles of relations, so there are three rules.

Path Consistency Rules

- Given $C_{x,y}, C_{x,z}, C_{y,z}$ replace $C_{x,y}$ by

$$C'_{x,y} = C_{x,y} \cap (C_{x,z} \cdot C_{y,z}^T)$$

- Given $C_{x,y}, C_{x,z}, C_{y,z}$ replace $C_{x,z}$ by

$$C'_{x,z} = C_{x,z} \cap (C_{x,y} \cdot C_{y,z})$$

- Given $C_{x,y}, C_{x,z}, C_{y,z}$ replace $C_{y,z}$ by

$$C'_{y,z} = C_{y,z} \cap (C_{x,y}^T \cdot C_{x,z})$$